

Score-linked practice intelligence from long-form violin training video

Alan N. Pham¹

¹AO Labs, Worcester, MA, USA

High-level music training produces dense evidence: daily practice videos, repeated repertoire, changing technical constraints, and periodic performance goals. Most of that evidence is difficult to use because it is stored as unstructured media rather than as a score-linked longitudinal record. We report Curtis, a live AO Labs system that converts messy long-form violin practice videos into structured, evidence-linked practice records using active practice detection, playable audio/video evidence windows, score alignment targets, heat-map targets, repertoire tracking, and passage-level performance observation. The checked live state on 8 May 2026 used the YouTube channel @nalalan as the primary source, indexed 211 public videos, marked 106 as practice candidates, marked 86 as long-form candidates, stored 46 media-sample index entries, retained 80 reviewed sections, and recorded 80 Curtis-focused findings. The uploaded-video ledger begins at violin 1, uploaded on 20 December 2025, and currently counts 42 violin-numbered or date-titled practice logs totaling 179 h 48 min of public session duration; this duration is not counted as active practice time. The live daily-record state contains 35 indexed practice days, 0 verified transcription records, 5 audio-evidence records, 13 violin-positive sample windows, and 31 withheld non-violin windows. Active violin practice time is counted only from media windows that are explicitly violin-positive; the current measured active total is 6 min 33 s, and all unverified room, setup, laptop, silence, or non-playing samples are withheld rather than counted as practice evidence. The primary interface has three top-level surfaces: a paper card, analyzed practice days, and repertoire. The analyzed-days evidence row separates indexed YouTube practice-archive duration, measured active violin time, and still-unmeasured archive duration so the interface does not hide the multi-hour corpus or overstate active hours. Each daily record now withholds embedded playback, paired audio, and active-time credit until the source window is violin-positive, and withholds sheet-music notation until score verification. Draft notation does not render as transcription; repeated single-pitch traces such as a spurious D D D

29 **D stream remain out of final notation and matching data. Repertoire entries are promoted only**
30 **from confirmed daily-record evidence and keep possible labels, user-rejected false labels, weak**
31 **evidence, and readiness scores inferred only from piece identification separate from confirmed**
32 **repertoire and judged playing quality. The system uses GPT-5.5 for vision and text review and**
33 **separate audio models for identification and verification, while preserving explicit limits: Curtis**
34 **admission cannot be predicted from current samples, weak audio/video evidence is marked**
35 **unjudged, whole-session transcription is not yet complete, and the system reports practice signals**
36 **rather than admissions probability. The contribution is an experimental method for making**
37 **practice visible, searchable, measurable, and tied to actual musical evidence.**

1 Introduction

Elite music preparation depends on repeated, deliberate practice and feedback over long time horizons (1, 2). The useful data are not only polished performances. They include ordinary practice sessions, failed takes, camera angle limitations, recurring technical problems, and changes in focus across days. A violinist recording daily practice can therefore produce a large longitudinal dataset, but the raw dataset does not automatically become an operating record.

Curtis addresses this problem as a live software system rather than as a portfolio page (3). The system indexes public practice media, separates uploaded video duration from detected active violin practice, records technical observations, and keeps evidence tied to clips and generated notation. It is designed for the Curtis Institute goal, but it does not claim to estimate an acceptance probability. The current role of the system is narrower: observe practice evidence, retain history, separate judged from unjudged evidence, and identify the specific passage-level blockers that can be tested in subsequent recordings.

The core pipeline is video → active practice detection → clips → music transcription → repeat grouping → score matching → heat map → repertoire update → passage-level observations. The output is intentionally small: analyzed practice days and an automatically updated repertoire list, both grounded in clips, notation, active playing time, and confidence limits.

The system is also part of AO Labs' broader progress architecture. In May 2026, AO Progress was introduced as a shared ledger for cross-application state and long-term changes across public apps, papers, CV artifacts, Curtis practice state, Imagineer state, and Relay state (4). Curtis supplies one stream in that ledger: timestamped media and review evidence for musical development.

2 System state

The production state uses the live backend at <https://curtis.aolabs.io/api/curtis/media-status>. The checked state on 8 May 2026 reported the values in Table 1. These are operating facts, not outcome claims.

2.1 Media indexing

The source inventory stores platform, title, publication time, duration, view count, candidate reasons, media kind, and blocker state. Dated long-form practice logs are treated as high-value longitudinal

Table 1. Curtis production state on 8 May 2026.

Signal	Current value
Primary source	YouTube channel @nalalan
Inventory	211 indexed videos
Practice candidates	106 videos
Long-form candidates	86 videos
Uploaded-video marker	violin 1, uploaded 20 December 2025
Title-confirmed uploaded-video ledger	42 videos, 179 h 48 min uploaded session duration
Active practice time	6 min 33 s counted only from violin-positive media windows; uploaded duration is reported separately
Unmeasured archive duration	179 h 41 min still awaiting full media segmentation
Media samples	46 indexed windows; 13 violin-positive; 31 withheld
Reviewed sections	80 sections
Curtis-focused findings	80 findings
Latest dated practice log	5-6-26
Review model	GPT-5.5 for vision/text, separate audio models for audio evidence and piece verification
Confirmed labels	5/1 Haydn Symphony No. 94 IV; 5/2 and 5/3 Wieniawski Scherzo-Tarantelle
Training anchors	3 source-confirmed practice sources
Score-reference targets	3 configured
Analyzed practice days	35 indexed days; 5 processed daily records; 5 audio-evidence records; 0 verified transcription records; 3 confirmed source labels with score snippets
Machine pitch events	Detected-active-section onset-aware pYIN plus onset-bounded spectral pitch rescue for fast passages, pitch-sanity filtering, repeated-pitch-collapse guarding, active-coverage fields, and fingerprints stored only when quality gates pass; unverified events remain paired audio evidence until score alignment is reliable
Specific observations	Rendered only when grounded in reliable score-linked evidence
Score-linked windows	5 target-linked windows; exact measure-level pitch/rhythm-to-score matching remains incomplete
Daily evidence surface	Local sample playback only for violin-positive windows, score heat-map target, notation withheld until score verification, and explicit whole-session transcription limit
Curtis-level completion	Not scored when the available evidence only identifies repertoire
Current focus	Audition-style setup consistency
Current constraint	Open strings only, mirror-check contact point
Boundary	Curtis admission cannot be predicted from current samples

records because the title contains a practice date and the duration preserves session-level evidence. The uploaded-video ledger starts at the `violin 1` marker and counts only violin-numbered or date-titled practice logs, so unrelated long public videos are retained in inventory without being silently counted as violin practice. This ledger is not the active-practice total. Active practice is measured only after media processing identifies windows where violin playing is actually occurring.

2.2 Review evidence

The review layer stores section-level findings with a dimension, evidence source, judgment, and practice constraint. The current dimensions include tone, intonation, shifts, time, articulation, and audition delivery. Findings can be strong, needing work, or unjudged. The unjudged state is part of the system design: if the sound is not available, the hand is obscured, the camera angle hides a transition, or still frames do not support a timing claim, the system should not convert weak evidence into a confident critique.

3 First readout

The first live readout is experimental and intentionally conservative. The strongest immediate value is that the system has already converted a large, timestamped practice archive into a queryable corpus with captured media samples, stored section reviews, and title-confirmed practice records. The home screen has three top-level surfaces: a paper card, analyzed practice days, and repertoire. The first evidence row now reports the YouTube practice archive first, then the active violin time that has actually been measured, then the archive duration still unmeasured; this prevents the 179 h 48 min public practice ledger from being visually collapsed into the much smaller processed-clip total. Each analyzed day groups same-day videos, preserves uploaded duration, and names confirmed or uncertain piece evidence without promoting guesses. Playable local media samples, paired audio/video evidence windows, and active-time credit are now withheld unless the source window is explicitly violin-positive. The owner-sync sampler now reuses locally cached captured media, uploads cached violin-positive windows before blind new captures, and probes deeper practice windows before spending extra passes on early setup segments; this added new violin-positive windows for 5/2, 5/3, 5/5, and 5/6. Sheet-music notation is also withheld until score verification, so audio evidence cannot appear as a failed transcription. Active scans with no violin-positive media remain pending evidence rather than

notation coverage. The repertoire surface is evidence-backed: an entry appears only when a daily record carries confirmed source evidence, and each entry keeps practice dates, clips, score targets, active time when measured, and current blocker limits. Generic etude/caprice evidence remains outside the piece-title field, and possible or uncorroborated repertoire labels remain checkable against the source window. These statements are not generic pedagogy; they are extracted from stored review findings, sample indexes, source corrections, public-domain score targets, pitch-quality checks, and the uploaded-video ledger.

Table 2. Current evidence classes and limits.

Class	Stored signal	Limit
Audio	Tone, intonation, shifts, time, articulation	Some audio reviews fail or cannot inspect the excerpt cleanly
Video	Setup, posture, bow lane, visible hand/bow evidence	Still frames can hide motion, sound, and transition accuracy
Metadata	Dates, duration, channel, view count, practice-candidate labels	Metadata cannot judge playing quality by itself
Progress plan	One focus, practice constraint, short session plan	The plan is a next experiment, not an outcome prediction

Current transcription limit. Machine note output now has two upstream gates. First, the media window must be classified as violin-positive before it can appear as playable evidence or count toward active time. Second, the note stream must be score-verified before it can render as sheet-music transcription. Unverified room, setup, laptop, silence, or non-playing samples remain withheld, and unverified pitch traces remain hidden support data rather than visible notation. A spectral-onset rescue path estimates pitch inside short onset-bounded violin segments when the pYIN path collapses, while the repeated-note trace guard prevents fast or complex playing from becoming a dominant repeated pitch, records the quality mode, and keeps unusable traces out of training data. Full-session score-linked transcription remains incomplete until every active playing section in the long-form practice videos is extracted, verified as violin-positive, and aligned to the score.

4 Discussion

Curtis Media Review is intentionally conservative. It does not rank the player against applicants, infer admissions probability, or convert every video into a claim. Instead, it keeps a running technical record. The sampler now treats violin presence as a hard evidence boundary: playable clips and active-time credit appear only after a media window passes the violin-positive gate, while notation waits for score verification. This makes the system useful even before it becomes musically sophisticated: it can preserve practice volume, surface recurring constraints, identify missing evidence, and make subsequent recordings easier to compare against earlier ones.

The main open problem is repertoire identification and evidence quality. YouTube metadata can identify candidate videos, and the current owner-sync path can provide direct audio/video samples, but unclear camera framing, missing score/title context, and short excerpts still prevent reliable piece naming. The backend therefore keeps generic labels out of the piece-title field, stores them as candidate evidence, requires source-window evidence for confirmed titles, and marks weak evidence explicitly. The current training ledger separates source-confirmed labels from independent audio-title matches, failed unverified machine pitch/rhythm windows, failed pitch-collapse windows, and score-aligned windows. It also stores reference targets for the three confirmed source labels across two works so later passes can ask for section-level matching instead of repeating broad title guesses. The sampler classifies candidate windows using active audio, violin-range voiced pitch, onset density, pitch-class diversity, spectral features, and minimum active-duration gates before a sample can become visible evidence. The transcription path converts sampled media to mono audio, finds audio-active playing sections, runs pitch tracking only on those sections, normalizes and harmonic-isolates each section for pitch tracking, estimates monophonic violin pitch with an onset-aware pYIN pitch track bounded to the violin range, detects onset boundaries so repeated attacks on the same pitch are not silently merged into one sustained note, runs a second spectral pitch estimate inside short onset-bounded segments when fast playing defeats pYIN, compares pitch diversity across candidate note streams, rejects out-of-range pitch artifacts, requires sustained pitch changes for pitch-track segmentation, removes very short low-confidence pitch blips, marks isolated octave-flip corrections as uncertain, rejects repeated-pitch collapse, estimates tempo, and builds pitch-class, interval, rhythm, and contour fingerprints only when quality gates pass. Daily records now carry transcription quality and active-coverage states, and the primary interface withholds media playback unless the source window is violin-positive and withholds machine notation until score verification. A stronger system should align those events against a symbolic or optical score representation, identify movement sections or measure ranges,

compare against reference recordings, and then judge articulation, timing, tone, and tempo against the aligned passage.

The broader experiment is whether a practice system can make musical development legible over time. A useful Curtis system should show whether tone, intonation, rhythm, articulation, shifts, memory, endurance, and audition delivery are improving across dated recordings. It should also preserve the source evidence behind each claim so that the dashboard does not become motivational text detached from the actual playing.

5 Materials and methods

Backend

The backend is a FastAPI service deployed on Railway. It exposes `https://curtis.aolabs.io/api/curtis/media-status` for compact state, `https://curtis.aolabs.io/api/curtis/ops-check` for fuller operational state, local stored sample media for embedded playback after the media is violin-positive, sample indexes for duplicate-window avoidance, and run endpoints for scans, media probing, analysis, and coaching. Persistent runtime state stores sources, inventory, review findings, media samples, analysis runs, and scan history.

Source scan

The default public source is `https://www.youtube.com/@nalalan`. The scanner queries public video metadata, classifies videos by title, duration, and candidate reasons, and stores inventory records. Dated practice logs and long-form videos are treated as practice candidates. The media probe now keeps one early baseline window but then probes deeper quarter, midpoint, five-eighths, three-quarter, and tail windows before spending additional passes on early setup probes, and each stored sample is keyed by video identifier plus window start so later passes can expand coverage without overwriting earlier evidence. Each sample is then classified by the violin-presence sampler before it can count as active practice evidence. The owner-media helper also checks its local cache of already captured browser samples, reuses classifier results when the file signature and classifier version match, and prioritizes unsynced violin-positive cached windows before capturing more blind windows. The practice-hours total uses a stricter title-confirmed ledger from violin 1 onward: violin-numbered videos and date-titled logs count toward the top-line session duration, while unrelated long videos

remain indexed but outside the practice-hours total. The scan does not itself inspect performance quality.

Model review

The model layer reviews sampled media sections where usable excerpts exist. GPT-5.5 is used for text and visual review, and separate audio models are configured for audio review and piece-title verification. Findings are sanitized so weak evidence terms such as obscured, not audible, or no clear evidence produce an unjudged state rather than an unsupported critique. Piece labels are also sanitized: coach reviews can store candidate evidence but cannot create confirmed repertoire titles, and the piece-identification layer must provide source-window evidence plus independent verification before it updates the current piece list. Same-video consensus can compare multiple sampled windows before accepting a title; for long date-stamped practice captures with later samples, setup or non-playing windows are excluded from confirmation, and confirmed active titles are versioned so older labels are rechecked when source rules change. User-rejected false labels are excluded from later passes only for the source where the correction applies, so later videos can still identify that repertoire when present. A source correction also invalidates older confirmed labels for that same source instead of promoting the next stale candidate, and repeated corrected false labels keep that source out of the confirmed repertoire list until a stronger acceptance path exists. Rejected proposed titles, top candidates, and verification titles are scrubbed from current API output, not only from confirmed piece records. The 5/1 forensic pass rejected prior Paganini, Bruch, Mendelssohn, Brahms, Sibelius, Tchaikovsky, Wieniawski, Saint-Saens, Ravel, Bazzini, Ernst, Carmen, Zigeunerweisen, Bach Partitas, Kreisler Praeludium and Allegro, Sarasate family labels, Ysaye Sonata No. 3 Ballade, Mozart K. 216, Glinka Ruslan and Lyudmila, Strauss Till Eulenspiegel, Strauss Don Juan, Smetana Bartered Bride, Prokofiev Classical Symphony, Dvorak Carnival Overture, Shostakovich Symphony No. 5, Ravel Bolero, the first broad orchestral candidate batch, and the later overture, dance, and pops-orchestral candidate batch for that source. Alan then supplied the ground-truth source label: Haydn Symphony No. 94, IV. Finale, Violin I part. The system now treats that label as human-confirmed for 5/1, stops injecting new seeded guesses for the same source, and keeps Curtis-level completion unscored until judged playing evidence supports a score. Alan also supplied the 5/2 source label, Wieniawski Scherzo-Tarantelle, Op. 16, and then confirmed that all violin footage on 5/3 was the same piece. The training state now records all three source labels as anchors and reports zero score-aligned windows until pitch/rhythm extraction has matched a source window to a score section. Transcribed samples

can now produce learned pitch/rhythm fingerprints only after their media window is violin-positive, and visible notation requires score verification. Future samples can receive a nearest-source hint when their extracted notes and rhythms resemble a prior confirmed source. Possible labels and model-only verified titles remain uncertain daily evidence unless a source label is confirmed, so future misses appear as missing evidence rather than as wrong repertoire. Piece-identification output does not create a Curtis-level readiness percentage by itself; readiness scoring requires judged playing evidence. Dated practice entries are keyed by the source-video date rather than the later processing time, so historical uploads cannot masquerade as today's practice state. The owner-browser sync uses the sample index to prioritize uncaptured public YouTube windows rather than re-uploading already reviewed excerpts, and the piece-identification pass can test multiple active moments inside each captured sample before withholding the title. The daily-record builder joins the strict violin 1 ledger to violin-positive media samples, hidden machine note/rhythm events, active-time evidence, clips, score packet support, repeated-fragment heat maps, and observation records. Repertoire entries are built only from confirmed daily-record evidence and keep progress percentages unassigned until clips, notation, and score alignment support them.

Progress ledger

AO Progress tracks Curtis as part of a larger multi-application state record. The first progress snapshot included Curtis home, Curtis media state, other AO Labs apps, Imagineer state, Imagineer paper, Relay state, and public project surfaces. Curtis contributes dated practice inventory, current focus, review status, and future trend signals.

6 Limitations

The current system cannot predict acceptance at Curtis. It also cannot replace a musician, teacher, jury, studio lesson, or audition process. It can be wrong when the source media are incomplete, when audio is unavailable, when video frames hide the physical action, or when model review fails. The useful claim is therefore limited: Curtis Media Review turns practice media into a persistent, source-bounded record that can support iteration over time.

References

227

1. K. A. Ericsson, R. T. Krampe, C. Tesch-Romer, The role of deliberate practice in the acquisition of expert performance. *Psychological Review* **100**, 363–406 (1993). 228
229

2. A. Williamon, Ed., *Musical Excellence: Strategies and Techniques to Enhance Performance* (Oxford University Press, Oxford, 2004). 230
231

3. A. N. Pham, *Curtis Media Review*, Live software system, curtis.aolabs.io, 2026. 232

4. A. N. Pham, *AO Progress Ledger*, Live longitudinal state system, progress.aolabs.io, 2026. 233